

Stanley 1913 has fueled the human experience since the iconic Stanley vacuum bottle revolutionized the way people enjoyed food and beverage. Today, our colorful and thoughtfully designed products including the popular Quencher™ series go beyond function, elevating everyday moments with style and empowering active lifestyles. We're a community of creators, builders and inventors who believe in creating sustainable products for a better life and world that minimize our impact on the planet. Learn more at [www.stanley1913.com](http://www.stanley1913.com).

## Position Overview

Stanley 1913 is seeking a Full Stack AI Engineer with experience in Generative AI, Multi Agent Systems. You will work across the entire stack—developing agentic pipelines using LangGraph, CrewAI, and Model Context Protocol (MCP) on Azure Cloud and Databricks. You will also perform rapid prototyping, POC development, and exploration of emerging AI technologies. Driven by curiosity and passion for innovative responsible applications of AI, you will ensure our business is able to grow nimbly and surprise and delight our customers. Your work will enable our builders to maintain a focus on their highest-impact work which will help our team and customers thrive.

## What You'll Do

- Build and orchestrate multi-agent systems using LangGraph, CrewAI, and MCP with well-defined agent logic, tools, and workflows.
- Implement retrieval-augmented generation (RAG) and agentic RAG pipelines with hybrid search, context compression, and memory modules.
- Deploy and fine-tune OpenAI and open-source LLMs (e.g., LLaMA, Mistral) using techniques like LoRA and QLoRA.
- Manage CI/CD workflows using GitHub Actions or Azure DevOps, and monitor model pipelines and health.
- Ensure the platform is scalable, highly performant, and secure
- Own projects from conception to deployment, monitoring, and iteration, driving end-to-end delivery.
- Engage product management and engineering to align project roadmaps, scope, and timelines.
- Stay ahead of with the latest advancements in LLMs and agentic AI, evaluating new approaches for practical applications.
- Foster a culture of experimentation and innovation, encouraging bold ideas and rapid delivery.

## Who You Are

- 5+ years of software engineering experience, with 1+ years focused on AI/ML systems.
- Experience with LangGraph, CrewAI, Model Context Protocol (MCP), RAG, and multi-agent orchestration,
- Experience in regression, classification, clustering, deep learning, NLP etc
- Experience with Python, Databricks, Spark, TensorFlow, PyTorch, and cloud platforms

- Experience with embedding models, document loaders, graph-based memory systems, and long-context handling.
- Strong Azure expertise, including infrastructure-as-code (Terraform), serverless computing, observability, and scaling.
- Familiarity with MLOps and AI application patterns (e.g., vector DBs, retrieval-augmented generation, MLFlow).
- Experience in evaluating models, and implementing guardrails
- Ability to push boundaries, experiment with innovative approaches, and drive continuous improvement in engineering practices.
- Experience in fast-paced environments where adaptability and rapid iteration are crucial.
- Frontend experience for building AI chat interfaces will be a plus

The base pay range for this position is for a successful candidate within the state listed. The successful candidate's actual pay will be based on multiple factors such as work location, job-related knowledge, skills, qualifications, and experience.

=====

#### What You Can Expect

As an AI Engineer, you will collaborate to design, implement, and optimize AI algorithms and software applications. You will ensure AI training, inference, deployment, and operation are functional, reliable, and scalable in your role.

#### About The Team

Zoom is seeking a highly passionate AI Software Engineer to join our AI infrastructure team. This AI Software Engineer plays a pivotal role in developing and enhancing AI-powered software solutions.

#### Responsibilities

- Collaborating with multidisciplinary teams to understand project requirements and objectives.
- Implementing AI solutions solves complex problems in computer vision, natural language processing and recommendation systems.
- Designing, developing, and maintaining AI algorithms, models, and software applications.
- Ensuring software reliability and quality by participating in code reviews, testing, and debugging.

- Documenting code, algorithms, and technical processes for internal and external reference.

#### What We're Looking For

- Possess a Bachelor's Degree in Computer Science, AI, Machine learning or a related field.
- Minimum 3+ years of industry experience or relevant work in a related field.
- Demonstrate deep understanding and skill in software engineering, AI, ASR or machine learning through experience.
- Be proficient in programming languages such as C++ programming; fluent Python programming; algorithm experience.
- Have experience with deep learning, neural networks, and data preprocessing.
- Have a solid comprehension of AI concepts, algorithms, and frameworks (e.g. TensorFlow, PyTorch).
- Be knowledgeable about software development best practices and version control systems (e.g. Git).

#### Minimum

#### Salary Range or On Target Earnings:

\$143,000.00

#### Maximum

\$312,800.00

In addition to the base salary and/or OTE listed Zoom has a Total Direct Compensation philosophy that takes into consideration; base salary, bonus and equity value.

Note: Starting pay will be based on a number of factors and commensurate with qualifications & experience.

We also have a location based compensation structure; there may be a different range for candidates in this and other locations.

#### Ways of Working

Our structured hybrid approach is centered around our offices and remote work environments. The work style of each role, Hybrid, Remote, or In-Person is indicated in the job description/posting.

## Benefits

As part of our award-winning workplace culture and commitment to delivering happiness, our benefits program offers a variety of perks, benefits, and options to help employees maintain their physical, mental, emotional, and financial health; support work-life balance; and contribute to their community in meaningful ways. Click [Learn](#) for more information.

## About Us

Zoomies help people stay connected so they can get more done together. We set out to build the best collaboration platform for the enterprise, and today help people communicate better with products like Zoom Contact Center, Zoom Phone, Zoom Events, Zoom Apps, Zoom Rooms, and Zoom Webinars.

We're problem-solvers, working at a fast pace to design solutions with our customers and users in mind. Find room to grow with opportunities to stretch your skills and advance your career in a collaborative, growth-focused environment.

## Our Commitment

At Zoom, we believe great work happens when people feel supported and empowered. We're committed to fair hiring practices that ensure every candidate is evaluated based on skills, experience, and potential. If you require an accommodation during the hiring process, let us know—we're here to support you at every step.

We welcome people of different backgrounds, experiences, abilities and perspectives including qualified applicants with arrest and conviction records and any qualified applicants requiring reasonable accommodations in accordance with the law.

If you need assistance navigating the interview process due to a medical disability, please submit an [Accommodations Request Form](#) and someone from our team will reach out soon. This form is solely for applicants who require an accommodation due to a qualifying medical disability.

Non-accommodation-related requests, such as application follow-ups or technical issues, will not be addressed.

Think of this opportunity as a marathon, not a sprint! We're building a strong team at Zoom, and we're looking for talented individuals to join us for the long haul. No need to rush your application – take your time to ensure it's a good fit for your career goals. We continuously review applications, so submit yours whenever you're ready to take the next step.

=====

radial is a Seattle-based startup enabling digital experiences at the speed of thought. We empower marketers and creatives to implement their ideas directly, with software that adapts over time. Our platform automates website and design system updates, large-scale migrations to new design systems, and continuous content optimization while adhering to company and product brands.

Backed by world class investors and a recent Series A raise of \$13 million—we're looking to scale our platform and expand our team. At Gradial, we operate with extreme ownership, bias towards action, and critical path planning. We tackle problems from first principles, question assumptions, and find creative solutions. If you want to take risks, work on groundbreaking technology, and see the direct impact of your work, Gradial is where you belong.

#### Role Overview

We're looking for an Applied AI Engineer to help build the next generation of agentic systems. This role combines deep software engineering skills with an experimental mindset around how AI can be practically applied to solve real user problems.

You'll collaborate across product, research, and engineering to design agents that reason, interact with environments and evolve through feedback. Ideal candidates bring curiosity, a love of ambiguous problems, and hands-on experience applying LLMs in production.

#### Responsibilities

- Build autonomous and semi-autonomous agent workflows that interact with browsers, codebases, and APIs to complete complex content operations.
- Evaluate model behavior in live environments and pioneer approaches for applying reinforcement learning to end-to-end agent systems.
- Design modular frameworks for tools, memory, and evaluators to enable continuous learning and improvement of agents.
- Develop data pipelines to collect, transform, and annotate real-world usage data for SFT and RL training.
- Design and build proprietary agent libraries, prompting strategies and benchmarking frameworks.
- Partner with customers and internal stakeholders to understand edge cases and align AI behaviors with user expectations.

## Requirements

- 4+ years of software engineering experience, with at least a year working directly on AI/ML-powered applications.
- Strong fluency in AI tooling (e.g. OpenAI APIs, Anthropic APIs).
- Practical understanding of LLMs, prompt engineering and agentic architectures.
- Experience building production systems that integrate external environments - code execution, browser automation or API orchestration.
- High degree of ownership and comfort in experimental, fast-moving environments.

## Desirable Attributes

- Familiarity with benchmarks like SWE-bench or evaluation strategies for tool-using agents.
- Exposure to customer-facing roles or user research for AI product development.
- Background in reinforcement learning or data-centric AI approaches.

The base salary range for this position is \$130,000 – \$220,000 annually. Final compensation will be determined based on factors such as experience, skills, and qualifications. In addition to base salary, this role may be eligible for performance-based bonuses and equity awards. Gradiant offers a comprehensive benefits package, including medical, dental & vision insurance, 401K retirement plan, paid time off, paid sick leave and other employee wellness programs.

## You'll thrive here if you...

- Learn quickly and actively seek out new challenges.
- Embrace AI as a core tool for problem-solving, creativity and scale.
- Show a strong work ethic, high ownership and bias toward action.
- Communicate clearly, directly and with curiosity.
- Thrive in fast-paced, hyper-growth environments where building better > maintaining status quo.

## AI Literacy & Interviewing Tools

As an AI-first company, we prioritize AI literacy as a core competency in our hiring decisions. We're excited by candidates who thoughtfully apply AI tools in their work, but during interviews we're focused on you. This is your opportunity to show how you think, communicate, and solve problems. Over-reliance on AI-generated responses during the interview process (especially when it obscures your own voice) will result in disqualification. We want to understand your unique perspective and how you approach challenges, both with and without AI.

## What We Offer

- Competitive salary and meaningful equity
- Comprehensive health, dental and vision coverage

- Fast-paced environment with flexibility and ownership
- Real impact, zero bureaucracy
- A front-row seat to building category-defining AI infrastructure

Gradiant is dedicated to creating an environment where diverse perspectives are valued and all team members can grow. We offer competitive compensation, equity, flexible work hours, comprehensive benefits, and a collaborative culture focused on learning and impact.

=====

## About the job

See yourself at Radian? We see you here too.

At Radian, we see you. For the person you are and the potential you hold. That's why we've embraced a new way of working that lets our people across the country be themselves, be their best and be their boldest. Because when each of us is truly seen, each of us gives our best – and at Radian, we'll give you our best right back.

Studies have shown that job seekers may hesitate to apply for jobs unless they meet every single qualification listed. We strive to see the potential in each applicant, so if you're excited about this role but your experience or education level doesn't align perfectly with every qualification in the job description, we encourage you to apply anyway. You may be just the right candidate for this or other roles.

### See Yourself as a Generative AI Engineer

As the Generative AI Engineer you are responsible for the development of innovative AI solutions. Your expertise in generative models, deep learning, and data analysis will be critical in creating intelligent and transformative AI applications. The resource will work closely with cross-functional teams to conceptualize, design, test, and deploy AI projects that drive innovation and provide value in the rapidly evolving field of artificial intelligence.

### See Your Primary Duties And Responsibilities

- Design and develop algorithms for generative models using deep learning techniques.

- Collaborate with cross-functional teams to integrate generative AI solutions into existing workflow systems.
- Research and stay up-to-date on the latest advancements in generative AI technologies and methodologies.
- Optimize and fine-tune generative models for performance and efficiency.
- Troubleshoot and resolve issues related to generative AI models and implementations.
- Create and maintain documentation for generative AI models and their applications.
- Communicate complex technical concepts and findings to non-technical stakeholders.
- Perform other duties as assigned or apparent.

See The Job Specifications

Your Basic Education And Prior Work-Related Experience

- Degree Requirement: Bachelor's Degree or Equivalent Experience
- Work Experience: 1-2 years of prior work-related experience

Additional Qualifications

- Strong Python skills; Effective prompt engineering techniques; familiarity with other languages is a plus.
- Proficiency in AWS Bedrock includes Model access and knowledge base implementations. Experience in Azure OpenAI.
- Hands-on with models for text (e.g., Claude, Openai gpt, Gemini, etc.), understand model fine tuning.
- Building efficient data pipelines and managing large datasets.
- Measuring model outputs with suitable metrics; awareness of bias and fairness.
- AI coding assistant (e.g. github copilot and etc), version control (Git), testing, and building scalable systems.
- Must have AWS serverless experience, Azure or GCP.
- Reading and implementing recent AI papers; experimentation.
- Documenting work and collaborating with teams.

See Your Location

Radian is committed to a flexible work environment for many of our roles. This is a *\*Work From Anywhere\** role meaning you have the flexibility to work from home (or another designated workspace that fits your needs).

This role provides additional flexibility should you want to work on-site at a Radian office. Explore our office locations [here](#) and let your Talent Acquisition Partner know you would be interested in working on-site.



- Work From Anywhere is subject to Radian's Alternative Work Policy and business needs.

#### See Why You Should Work With Us

- Competitive Compensation: anticipated base salary from \$78,700 to \$112,400 based on skills and experience. This position is eligible to participate in an annual incentive program.
- Rest and Relaxation. This role is eligible for 20 days of paid time off annually, which is prorated in the year of hire based on hire date. In addition, based on your hire date, you will be eligible for 9 paid holidays + 2 floating holidays. Parental leave is also offered as an opportunity for all new parents to embrace this exciting change in their lives.
- Our Company Makes an Impact. We've been recognized by multiple organizations like Bloomberg's [Gender-Equality Index](#), HousingWire's [Tech 100](#), and The Forum of Executive Women's [Champion of Board Diversity](#). Radian has also pledged to PwC's [CEO Action for Diversity & Inclusion](#) commitment.
- Comprehensive Health Benefits. Multiple medical plan choices, including HSA and FSA options, dental, vision, and basic life insurance.
- Prepare for your Future. 401(k) with a top of market company match (did we mention the company match is immediately vested?!) and an opportunity to participate in Radian's Employee Stock Purchase Plan (ESPP).
- Homebuyer Perks. Our Homebuyer Perks program helps employees navigate the home searching, buying, selling, and refinancing processes and provides valuable financial benefits to encourage, enable, and support home ownership.
- Additional Benefits. To learn more about our benefits offerings, visit our [Benefits Page](#).

*The application period for the job is estimated to be 20 days from the job posting date. However, this timeline may be shortened or extended depending on business needs and the availability of qualified candidates.*

Radian will consider for employment qualified applicants with arrest or conviction records in a manner consistent with the requirements of the law, including any applicable fair chance law.

#### See More About Radian

Radian is a fintech servicing the mortgage and real estate services industry. As a team, we pride ourselves on seeing the potential of every person, every idea and every day.

Seeing each other at Radian goes far beyond our open, flexible culture. It means seeing our people's potential – and creating inspiring career paths that help them get there. Or seeing new pathways and innovating for the future of our industry. It means seeing each other for all that we are. And it means seeing our purpose as one that extends beyond the bottom line – having an impact on communities across the country to help more people achieve the American Dream of homeownership.

We hope you'll see yourself at Radian. See more about us at [Radian.com](#).

## Defining Roles for Radian's Future

Understanding the qualities and characteristics that define a Leader and an Employee is important to building our future-fit workforce. Radian's future is only as bright as its people. For that reason, our People Plan includes profiles to support the qualities and characteristics that each [Leader](#) as well as each [Employee](#) should embody upon hire or via development.

=====

## About the job

CSQ426R189

The AI Forward Deployed Engineering (AI FDE) team is a highly specialized customer-facing AI team at Databricks. We deliver professional services engagements to help our customers build and productionize first-of-its-kind AI applications. We work cross-functionally to shape long-term strategic priorities and initiatives alongside engineering, product, and developer relations, as well as support internal subject matter expert (SME) teams. We view our team as an ensemble: we look for individuals with strong, unique specializations to improve the overall strength of the team. This team is the right fit for you if you love working with customers, teammates, and fueling your curiosity for the latest trends in GenAI, LLMOps, and ML more broadly. This role can be remote.

## The Impact You Will Have

- Develop cutting-edge GenAI solutions, incorporating the latest techniques from our [Mosaic AI Research](#) to solve customer problems
- Own production rollouts of consumer and internally facing GenAI applications
- Serve as a trusted technical advisor to customers across a variety of domains
- Present at conferences such as [Data + AI Summit](#), recognized as a thought leader internally and externally
- Collaborate cross-functionally with the product and engineering teams to influence priorities and shape the product roadmap

## What We Look For

- Experience building GenAI applications, including RAG, multi-agent systems, Text2SQL, fine-tuning, etc., with tools such as HuggingFace, LangChain, and DSPy
- Expertise in deploying production-grade GenAI applications, including evaluation and optimizations
- Extensive years of hands-on industry data science experience, leveraging common machine learning and data science tools (i.e., pandas, scikit-learn, PyTorch, etc.)

- Experience building production-grade machine learning deployments on AWS, Azure, or GCP
- Graduate degree in a quantitative discipline (Computer Science, Engineering, Statistics, Operations Research, etc.) or equivalent practical experience
- Experience communicating and/or teaching technical concepts to non-technical and technical audiences alike
- Passion for collaboration, life-long learning, and driving business value through AI
- [Preferred] Experience using the Databricks Intelligence Platform and Apache Spark™ to process large-scale distributed datasets

## Pay Range Transparency

Databricks is committed to fair and equitable compensation practices. The pay range(s) for this role is listed below and represents the expected base salary range for non-commissionable roles or on-target earnings for commissionable roles. Actual compensation packages are based on several factors that are unique to each candidate, including but not limited to job-related skills, depth of experience, relevant certifications and training, and specific work location. Based on the factors above, Databricks anticipated utilizing the full width of the range. The total compensation package for this position may also include eligibility for annual performance bonus, equity, and the benefits listed above. For more information regarding which range your location is in visit our page [here](#).

### Zone 1 Pay Range

\$161,280—\$225,792 USD

### Zone 2 Pay Range

\$161,280—\$225,792 USD

### Zone 3 Pay Range

\$161,280—\$225,792 USD

### Zone 4 Pay Range

\$161,280—\$225,792 USD

## About Databricks

Databricks is the data and AI company. More than 10,000 organizations worldwide — including Comcast, Condé Nast, Grammarly, and over 50% of the Fortune 500 — rely on the Databricks Data Intelligence Platform to unify and democratize data, analytics and AI. Databricks is headquartered in San Francisco, with offices around the globe and was founded by the original creators of Lakehouse, Apache Spark™, Delta Lake and MLflow. To learn more, follow Databricks on [Twitter](#), [LinkedIn](#) and [Facebook](#).

=====

## About the job

### Company Overview:

Centific expertly engineers platforms and curates multimodal, multilingual data to empower the 'Magnificent Seven' and enterprise clients with safe, scalable Artificial intelligence (AI) deployment. Our team includes over 150 PhDs and data scientists, along with more than 4,000 AI practitioners and engineers. We leverage an integrated ecosystem comprised of industry-leading partnerships, and 1.8 million vertical domain experts across 230 locales, to create high-quality pre-trained datasets, fine-tuned industry-specific Large Language Models (LLMs), and Retrieval-Augmented Generation (RAG) pipelines supported by vector databases. Our innovations can reduce Generative Artificial Intelligence (Gen AI) costs by up to 80% and bring Gen AI solutions to market 50% faster. Our mission is to bridge the gap between AI creators and industry leaders by bringing best practices in Generative AI to unicorn innovators and enterprise customers. We aim to help these organizations unlock significant business value by leveraging Generative AI at scale, ensuring they stay at the forefront of technological advancement and maintain a competitive edge in their respective markets.

Website - <https://www.centific.com/>

Title: AI/ML Engineer – LLMs, Agents, Simulation

Location: Bellevue, WA (5 days a week onsite)

Stack: Python (Primary), C# (Optional), LLMs, Optimization & Simulation

We're applying frontier AI and LLM technology to domains where outcomes are measured in lives saved, logistics won, and missions completed. As an AI/ML Engineer, you will lead the integration of large language models into agentic workflows, simulation platforms, and real-time decision engines. If you've ever wanted to work on the bleeding edge of LLM integration, agent design, and AI reasoning at scale, this is your opportunity. You'll work with a world-class team where deep AI meets domain expertise—and where your work will directly influence national security.

### What You'll Do:

- Build and optimize agentic workflows powered by LLMs
- Drive prompt engineering and architecture for embedded reasoning
- Partner with simulation and backend teams to train and deploy models in real-world scenarios
- Contribute to novel approaches in AI-guided decision making in contested environments

### Ideal Candidate:

- 3+ years of experience in AI/ML or applied data science
- Expertise in Python and modern LLMs (e.g., OpenAI, Anthropic, open-source stacks)
- Familiarity with optimization, simulation modeling, or autonomous agents

- Bonus: experience with secure systems, C#, or decision modeling

You're not just writing code—you're helping defend the future, using AI as a force multiplier for good.

Benefits offered - comprehensive healthcare, dental & vision, 401k plan, PTO, etc.

Centific is an equal-opportunity employer. All qualified applicants will receive consideration for employment without regard to race, color, religion, national origin, ancestry, citizenship status, age, mental or physical disability, medical condition, sex (including pregnancy), gender identity or expression, sexual orientation, marital status, familial status, veteran status, or any other characteristic protected by applicable law. We consider qualified applicants regardless of criminal histories, consistent with legal requirements.

=====

## About the job

### About Ramp

At Ramp, we're rethinking how modern finance teams function in the age of AI. We believe AI isn't just the next big wave. It's the new foundation for how business gets done. We're investing in that future — and in the people bold enough to build it.

Ramp is a financial operations platform designed to save companies time and money. Our all-in-one solution combines payments, corporate cards, vendor management, procurement, travel booking, and automated bookkeeping with built-in intelligence to maximize the impact of every dollar and hour spent. But we're not just building features powered by AI. We're building a platform where agents can chase receipts, close books, flag risks, and surface insights. This enables teams to reclaim their time and reinvest in what matters.

More than 40,000 businesses, from family-owned farms to space startups, have saved \$10B and 27.5M hours with Ramp. Founded in 2019, Ramp powers the fastest-growing corporate card and bill payment platform in America, and enables over \$80 billion in purchases each year.

Ramp's investors include Thrive Capital, Sands Capital, General Catalyst, Founders Fund, Khosla Ventures, Sequoia Capital, Greylock, and Redpoint, among others, in addition to 100+ angel investors who have been founders or executives of leading companies.

Ramp has been named to Fast Company's [Most Innovative Companies](#) list and LinkedIn's [Top U.S.](#)

[Startups](#) for more than 3 years, as well as the Forbes [Cloud 100](#), CNBC [Disruptor 50](#), and TIME Magazine's 100 [Most Influential Companies](#).

## About The Role

We are fundamentally an engineering-driven company, and as we enter our next chapter, we're investing deeply in AI-native products to reshape how software empowers businesses. We believe in building full-stack, vertically integrated solutions that redefine financial workflows through automation, agents, and intelligent interfaces.

We're looking for an experienced AI Product Engineer to join our early efforts in reimagining core Ramp experiences through an AI-first lens. This role is ideal for a builder who thrives at the intersection of AI capabilities and product experience — someone excited to craft intelligent, user-facing features from the ground up and push the boundaries of what's possible with LLMs and agentic workflows.

You'll work alongside product, engineering, and design to build novel customer-facing products and AI-powered internal tooling. This is a highly cross-functional role with high ownership, a fast feedback loop, and deep impact on how Ramp's customers work.

Check out our [Engineering Blog](#) for more on our tech stack, mission, and values!

## What You'll Do

- Architect and ship new AI-native product experiences across Ramp, from 0→1 prototypes to scaled production systems
- Build LLM-powered agents that autonomously reason through complex workflows
- Contribute to our evolving AI product stack: prompt orchestration, context engineering, evals, embeddings, and retrieval systems
- Develop full-stack product surfaces (backend & frontend) that incorporate AI into user experiences
- Collaborate closely with AI platform, infra, and design to push the capabilities of our internal and external tooling
- Stay on the bleeding edge of the AI ecosystem (tools like Claude Code, Cursor, OpenAI Agent SDK, MCPs, etc.) and identify where and how we can adopt emerging techniques

- Influence Ramp's broader AI product strategy and technical vision

#### What You'll Need

- 2+ years of software engineering experience, including time spent building production-grade full-stack systems
- Demonstrated experience building AI-powered product features
- Fluency with modern LLM platforms and AI-assisted engineering tools (e.g., Cursor, Claude Code), with a bias toward rapid experimentation and a working knowledge of prompt engineering, retrieval, and orchestration
- Strong backend fundamentals (Python, Postgres, etc.)

#### Nice to Haves

- Experience building or scaling AI-native workflows, agents, or dev tooling (e.g., Modal, Turbopuffer, OpenAI Agent SDK)
- Familiarity with the AI development stack (RAG, vector databases, LLM observability, eval pipelines)
- Familiarity with the best AI products and the UX paradigms
- A deep interest in being an early adopter of the best AI products and bringing their learnings into Ramp
- Frontend experience using React/Next.js
- Exposure to Ramp's backend technologies or interest in learning them: Flask, SQLAlchemy, Postgres, Celery, and Temporal
- Prior startup experience or track record of building 0→1 product experiences in fast-paced environments

#### Benefits (for U.S.-based Full-time Employees)

- 100% medical, dental & vision insurance coverage for you
  - Partially covered for your dependents
  - One Medical annual membership
- 401k (including employer match on contributions made while employed by Ramp)
- Flexible PTO
- Fertility HRA (up to \$5,000 per year)
- WFH stipend to support your home office needs
- Wellness stipend
- Parental Leave
- Relocation support to NYC or SF (as needed)
- Pet insurance

=====

## About the job

## About The Role

Join Shopify's innovative team as we develop an AI Personal Shopper to transform the online shopping experience. Leveraging cutting-edge AI, including Large Language Models (LLM) and advanced machine learning algorithms, you'll play a pivotal role in delivering personalized recommendations and insightful suggestions tailored to individual preferences. Our goal is to redefine e-commerce by creating a concierge service that enhances how customers interact with Shop and Storefronts. As a Machine Learning Engineering (MLE) lead or individual contributor, you'll be at the forefront of implementing AI systems at scale, directly empowering merchants and creating tangible solutions with real-world impact.

### Key Responsibilities:

- Develop and deploy Generative AI, natural language processing, and machine learning models.
- Design and produce scalable AI/ML system architectures.
- Implement data pipelines for fine-tuning LLMs.
- Solve high-impact data problems, delivering business impact through data and machine learning products.
- Prioritize and communicate effectively with both technical and non-technical audiences.

### Qualifications:

- Mastery in building data products using generative AI, RLHF, and fine-tuning LLMs.
- End-to-end experience in training, evaluating, testing, and deploying machine learning products at scale.
- Experience in building data pipelines and driving ETL design decisions using disparate data sources.
- Proficiency in Python, shell scripting, streaming and batch data pipelines, vector databases, DBT, BigQuery, BigTable, or equivalent, and orchestration tools.
- Experience with running machine learning in parallel environments (e.g., distributed clusters, GPU optimization).
- This role may require on-call work.

*At Shopify, we pride ourselves on moving quickly—not just in shipping, but in our hiring process as well. If you're ready to apply, please be prepared to interview with us within the week. Our goal is to complete*



*the entire interview loop within 30 days. You will be expected to complete a pair programming interview, using your own IDE.*

*This role may require on-call work.*

Ready to redefine e-commerce through AI innovation? Join the team that's making commerce better for everyone.

## About Shopify

Opportunity is not evenly distributed. Shopify puts independence within reach for anyone with a dream to start a business. We propel entrepreneurs and enterprises to scale the heights of their potential. Since 2006, we've grown to over 8,300 employees and generated over \$1 trillion in sales for millions of merchants in 175 countries.

This is life-defining work that directly impacts people's lives as much as it transforms your own. This is putting the power of the few in the hands of the many, is a future with more voices rather than fewer, and is creating more choices instead of an elite option.

## About You

Moving at our pace brings a lot of change, complexity, and ambiguity—and a little bit of chaos. Shopifolk thrive on that and are comfortable being uncomfortable. That means Shopify is not the right place for everyone.

Before you apply, consider if you can:

- Care deeply about what you do and about making commerce better for everyone
- Excel by seeking professional and personal hypergrowth
- Keep up with an unrelenting pace (the week, not the quarter)
- Be resilient and resourceful in face of ambiguity and thrive on (rather than endure) change
- Bring critical thought and opinion
- Use AI tools reflexively as part of your fundamental workflow
- Embrace differences and disagreement to get shit done and move forward
- Work digital-first for your daily work

=====

## About the job

Hudson River Trading (HRT) is hiring an AI Researcher to join the HAIL team. HAIL (HRT AI Labs) is the team at HRT responsible for developing and maintaining our most powerful models, which are used by our trading teams to drive a significant fraction of our trading. We are building and deploying "foundation models for markets", that ingest and train on vast amounts of market data, to make predictions about future market state. We are seeking experienced AI researchers to join our team to

accelerate our efforts.

Researchers have great independence to pursue the research directions they think would be most impactful as part of a small focused team with minimal bureaucracy. They are enabled by state-of-the-art research clusters with very high GPU-to-researcher ratios, and supported by excellent engineering, hardware, and systems teams to realize their vision. Your work will be directly, clearly, highly impactful on the business, and it will be challenging: this is a field with no easy or obvious solutions. You'll be responsible for improving every part of our models: from featurization of data, to architecture design, to training dynamics, to how trading decisions are made.

### Qualifications

- Must have two or more years work experience in deep learning research, for any domain: robotics, biology, chemistry, physics, audio, video, etc.
- Experience translating methods between areas of application is highly valued.
- LLM experience is valuable, but not necessary.
- Strong engineering skills valued, especially any of: GPU kernel development, lower-level Pytorch/JAX development, intermediate C++, FPGA/ASIC experience, HPC experience.

The estimated base salary range for this position is \$200,000 to \$300,000 per year, based on job-related skills and experience. This role will also be eligible for discretionary performance-based bonuses and a competitive benefits package which includes medical, dental, vision, basic life insurance, and enrollment in our company's 401k plan. Employees will also receive 20 vacation days and 10 paid holidays annually, as well as sick and parental leave.

### Culture

Hudson River Trading (HRT) brings a scientific approach to trading financial products. We have built one of the world's most sophisticated computing environments for research and development. Our researchers are at the forefront of innovation in the world of algorithmic trading.

At HRT we welcome a variety of expertise: mathematics and computer science, physics and engineering, media and tech. We're a community of self-starters who are motivated by the excitement of being at the cutting edge of automation in every part of our organization—from trading, to business operations, to recruiting and beyond. We value openness and transparency, and celebrate great ideas from HRT veterans and new hires alike. At HRT we're friends and colleagues –

whether we are sharing a meal, playing the latest board game, or writing elegant code. We embrace a culture of togetherness that extends far beyond the walls of our office.

Feel like you belong at HRT? Our goal is to find the best people and bring them together to do great work in a place where everyone is valued. HRT is proud of our diverse staff; we have offices all over the globe and benefit from our varied and unique perspectives. HRT is an equal opportunity employer; so whoever you are we'd love to get to know you.

=====

## **About the job**

On this journey for over 12 years, Qventus is leading the transformation of healthcare. We enable hospitals to focus on what matters most: patient care. Our innovative solutions harness the power of machine learning, generative AI, and behavioral science to deliver exceptional outcomes and empower care teams to anticipate and resolve issues before they arise.

Our success in rapid scale across the globe is backed by some of the world's leading investors. At Qventus, you will have the opportunity to work with an exceptional, mission-driven team across the globe, and the ability to directly impact the lives of patients. We're inspired to work with healthcare leaders on our founding vision and unlock world-class medicine through world-class operations.

### About The Role

We're looking for a Solutions Architect, Conversational AI & Prompt Engineering to lead the design and optimization of intelligent, scalable, and modular chatbot systems. You'll own the end-to-end architecture of conversational flows and collaborate across data science, engineering, and product teams to build high-quality conversational AI experiences.

This role blends deep technical expertise in prompt engineering, LLM integration, and conversational design, with strong system thinking and collaboration skills. You'll also help structure prompt logic across broader AI-driven product experiences.

### Key Responsibilities

#### Conversational Logic & Prototyping

- Design and implement dynamic AI conversation flows that handle multi-turn interactions, interruptions, and alternate pathways.
- Build and manage systems for dialogue state, intent recognition, and entity extraction.
- Architect reusable, scalable chatbot logic using a modular design approach.

#### Technical Implementation

- Leverage LLM-based AI (GPT, Claude, Gemini, LLaMA, etc.) to build advanced conversational systems.
- Apply conversation modeling techniques like state machines, decision trees, and graph-based flows.
- Use frameworks such as Dialogflow, Rasa, Amazon Lex, Langchain, or similar tools.
- Collaborate with backend teams to integrate bots with APIs, CRMs, and enterprise systems.
- Enhance bot functionality with custom scripts, API calls, and external system integrations.

#### Optimization, Testing & Maintenance

- Continuously tune and improve intent recognition and response quality using real user data.
- Monitor and resolve conversation breakdowns, mismatches, and fallback issues.
- Develop and run unit, regression, and A/B tests to ensure conversational performance.
- Implement real-time monitoring, analytics, and logging to drive performance insights.

#### Collaboration & Documentation

- Partner with product teams to align conversational UX with user and business goals.
- Work closely with software engineers and data scientists to refine LLM integration and behavior.
- Document conversation logic, prompt structures, and decision workflows for scalability and maintainability.

#### It's a Plus If You Have

- Experience with voice-enabled conversational AI systems.
- Knowledge of vector databases, embeddings, and RAG (Retrieval Augmented Generation).
- Experience applying knowledge graphs in chatbot workflows.
- Familiarity with ML techniques for intent classification and dialog management.
- Prior experience building chatbots in healthcare or other regulated industries.

Compensation for this role is based on market data and takes into account a variety of factors, including location, skills, qualifications, and prior relevant experience. Salary is just one part of the total rewards package at Qventus. We also offer a range of benefits and perks, including Open Paid

Time Off, paid parental leave, professional development, wellness and technology stipends, a generous employee referral bonus, and employee stock option awards.

=====

## About the job

Cresta is on a mission to turn every customer conversation into a competitive advantage by unlocking the true potential of the contact center. Our platform combines the best of AI and human intelligence to help contact centers discover customer insights and behavioral best practices, automate conversations and inefficient processes, and empower every team member to work smarter and faster. Born from the prestigious Stanford AI lab, Cresta's co-founder and chairman is [Sebastian Thrun](#), the genius behind Google X, Waymo, Udacity, and more. Our leadership also includes CEO, [Ping Wu](#), the co-founder of Google Contact Center AI and Vertex AI platform, [and co-founder, Tim Shi](#), an early member of Open AI.

Join us on this thrilling journey to revolutionize the workforce with AI. The future of work is here, and it's at Cresta.

### About The Role

At Cresta, the AI Agent team is on a mission to create state-of-the-art AI Agents that solve practical problems for our customers. We are focused on leveraging the latest technologies in Large Language Models (LLMs) and AI Agent systems, while ensuring that the solutions we develop are cost-effective, secure, and reliable.

As an AI Agent Engineer, you'll be at the forefront of deploying AI agents that address real-world challenges. In this role, you will work closely with customers as well as our software and machine learning engineers, ensuring high-impact AI Agent deployments and contributing to the continuous improvement of our core AI platform. You'll develop intelligent AI agents, integrate them seamlessly with external systems and offer hands-on technical expertise to ensure successful deployments.. This position requires strong engineering skills, adaptability, and customer engagement. If you are self-driven, analytical, and eager to leverage AI in practical applications, this role is for you.

### What You'll Do

- **AI Agent Development:** Develop, configure, deploy, and optimize AI agents using Cresta's AI platform and tools.
- **System Integration:** Build AI agent integrations with external systems (APIs, databases, CRMs, etc.) to ensure seamless workflow integration.
- **Performance Optimization:** Optimize AI agent performance (e.g. fine-tune prompts and configurations) and troubleshoot issues in complex enterprise environments.
- **Requirements Gathering:** Collaborate with customers and internal stakeholders to gather technical requirements and translate business needs into AI Agent solutions.
- **Demos & Iteration:** Conduct interactive demos and present compelling proof-of-concepts to prospective customers, proactively gather feedback, and iteratively refine solutions to meet objectives.
- **Project Coordination:** Define project milestones, create implementation plans, and coordinate execution with internal teams to ensure on-time delivery. Provide a tight feedback loop to our product and engineering teams – identifying gaps, building custom tooling, and influencing the roadmap through real-world deployment learnings.
- **Partner Cross-Functionally:** Collaborate with PMs to define agent goals, iterate rapidly based on customer feedback, and shape product capabilities that maximize customer ROI.
- **Technical Guidance:** Serve as a trusted technical advisor for the customer, guiding best practices for AI agent adoption and usage. Provide technical guidance on AI agent best practices, including architecture design, security considerations, and scalability planning.

## What We're Looking For

- **Education:** Bachelor's or Master's degree in Computer Science, Engineering, or a related field.
- **Experience:** 3+ years of experience in software development, AI/ML engineering, or system integration.
- **Programming Skills:** Proficiency in Python and Golang, with the ability to write clean, efficient code.
- **AI/ML Knowledge:** Familiarity with AI/ML concepts. Hands-on experience with large language models (LLMs), and prompt engineering techniques are strongly preferred.
- **AI Agent Frameworks:** Strong understanding of general AI agent frameworks, function calling, and retrieval-augmented generation (RAG). Hands-on experience of building such a system is strongly preferred.
- **Cloud & DevOps:** Experience with cloud platforms (AWS, GCP, or Azure) and DevOps practices (CI/CD, containerization, monitoring).
- **Integration Expertise:** Hands-on experience with integrating systems via APIs, webhooks, and data pipelines.
- **Communication:** Excellent communication and project management skills.
- **Analytical Approach:** Ability to use data-driven decision-making, including A/B testing and performance monitoring, to refine solutions.
- **Collaborative Builder:** You thrive in cross-functional environments, working hand-in-hand with PMs and engineers to turn real customer problems into scalable AI solutions.

## Perks & Benefits

We offer a comprehensive and people-first benefits package to support you at work and in life:

- Comprehensive medical, dental, and vision coverage with plans to fit you and your family
- Flexible PTO to take the time you need, when you need it
- Paid parental leave for all new parents welcoming a new child
- Retirement savings plan to help you plan for the future
- Remote work setup budget to help you create a productive home office
- Monthly wellness and communication stipend to keep you connected and balanced
- In-office meal program and commuter benefits provided for onsite employees

### Compensation At Cresta

Cresta's approach to compensation is simple: recognize impact, reward excellence, and invest in our people. We offer competitive, location-based pay that reflects the market and what each individual brings to the table.

The posted base salary range represents what we expect to pay for this role in a given location. Final offers are shaped by factors like experience, skills, education, and geography. In addition to base pay, total compensation includes equity and a comprehensive benefits package for you and your family.

Salary Range: \$150,000–\$250,000 + Offers Equity

We have noticed a rise in recruiting impersonations across the industry, where scammers attempt to access candidates' personal and financial information through fake interviews and offers. All Cresta recruiting email communications will always come from the @cresta.ai domain. Any outreach claiming to be from Cresta via other sources should be ignored. If you are uncertain whether you have been contacted by an official Cresta employee, reach out to [recruiting@cresta.ai](mailto:recruiting@cresta.ai)

=====

## About the job

Popl is hiring!

Join us on our mission of redefining how companies connect with their future customers.

About Popl

Popl is the leading platform for in-person lead capture and digital business cards. We help modern sales and marketing teams ditch outdated tools, like paper business cards and clunky badge scanners, for a smarter, faster way to capture and convert leads.

With our AI-powered mobile and web platform, teams can instantly collect contact info from event badges, business cards, and QR codes - then sync it all directly to their CRM. No manual entry. No lost leads. Just streamlined networking and faster follow-ups.

We're trusted by forward-thinking teams at Salesforce, CapGemini, Prudential, Hilti, the Brooklyn Nets, and thousands more. Whether it's a Fortune 500 or a fast-growing startup, companies use Popl to make every connection count.

Why join us?

We're an ambitious, creative, and customer-obsessed team - energized by solving big problems and building tools that bring human connection back to the forefront of business. We move fast, we love feedback, and we use our own product every single day.

You'll be joining a company that's:

- ✓ Backed by Y Combinator
- ✓ Growing fast (and profitably)
- ✓ Used and loved by millions of companies and professionals globally
- ✓ Built by a passionate of founders focused on people and culture
- ✓ Creating magical, "why didn't this exist before?" experiences

Who we're looking for:

We're always on the lookout for ambitious, driven people across roles in product, marketing, engineering, customer success, and sales.

You might be a great fit if you:

- Want to help build a category-defining product
- Love fast-moving startup environments
- Care deeply about the customer experience
- Believe in the power of face-to-face connection (and great software)
- Are excited to work with a high-energy, no-ego team

Let's connect!

Even if there's not a role that fits you perfectly, we'd still love to hear from you.

Check out [popl.co](https://popl.co) to learn more, or drop us a message.

=====

## About the job

Nuance Labs is an early-stage deep tech startup. We're building the first real-time human foundation model — unifying text, speech, and vision — to make AI socially and emotionally intelligent. Imagine an AI that can understand a quirked eyebrow, a shift in tone, or a hesitant pause, and respond in a way that feels truly human.

Key Facts



- \$10M seed round backed by Accel, South Park Commons, Lightspeed, and angel investors including the former Chief Product Officer of Synthesia.
- In-person collaboration, 5 days a week, Seattle HQ (with plan to expand to San Francisco)

This is for you, if you:

- Have at least 1 year of industry experience designing and building distributed ML pipelines, developer tooling, visualization systems, or other infrastructure that supports large-scale machine learning.
- Can partner closely with researchers to turn vague needs (and sometimes half-baked ideas) into robust, usable tools – often anticipating problems and solutions before they’re spelled out.
- Move with urgency and iterate fast.
- Write code that’s clean enough your future self will thank you for
- Play well with other brilliant minds from different domains

What you’ll be building

The first human foundation model that operates across text, speech, facial expression, and body language in real time. This unified model:

- Understands fine-grained human signals – from a quirked eyebrow to a subtle change in voice – and infers meaning in context
- Generates lifelike, responsive avatars whose expressions, gestures, and tone evolve frame-by-frame to deliver genuine responses

The landscape is ripe for innovation. While voice AI systems have made great strides in capturing prosody, and avatar platforms can generate compelling visuals, existing solutions remain fragmented. Real-time, multimodal interaction – where voice, facial expression, and contextual perception converge – is still an unsolved problem. This role offers the rare opportunity to shape foundational technology in a space where the boundaries are still being defined.

Why now?

As IQ in AI becomes a commodity, EQ is emerging as the key differentiator.

- Market - Social interaction is increasingly virtual, across work, healthcare, education, and entertainment. Millions already engage daily with AI companions in text. The leap to rich visual and emotional interaction is inevitable.
- Technology - Breakthroughs in expressive avatars, streaming speech, and real-time video generation now make this vision achievable at human speed.

Why this team?

We’re former Apple research scientists who’ve spent years advancing AI avatar and audio-visual generation – publishing at top conferences and shipping ultra-low-latency ML products to millions. We combine frontier research with the ruthless engineering needed for consumer-grade, real-time systems.

- Fangchang Ma: CEO. MIT PhD in Robotics and ML. Previously Research Manager at Apple, with experience at DJI. Published in top AI conferences with 2400+ citations.

- Edward Zhang: CTO. UW PhD in Computer Graphics. Previously Senior research scientist at Apple, with experience at Google, Microsoft.
- Karren Yang: Founding Research Scientist. MIT PhD. Previously senior research scientist at Apple, with experience at Niantic Labs, Meta Reality Labs, Bosch Center for AI, and Adobe Research. Published in top AI conferences with 1800+ citations.

We're a small, fast-moving research team with an exceedingly high bar — bringing on only the very best talent. Every member has massive ownership, deep trust, and the opportunity to shape both the technology and the company from the ground up.

#### Cultural Values

- Doing right by people — integrity and respect aren't negotiable.
- Transparency — we communicate openly and share context so everyone can make informed decisions.
- Relentless speed — we bias toward action, iterate fast, and learn quickly.
- In-person collaboration: the best ideas and strongest teams happen face-to-face. Being together fuels our energy, accelerates problem-solving, and sparks creativity.

#### How to Apply

Email [careers@nuancelabs.ai](mailto:careers@nuancelabs.ai) with your CV and a short note on why this role excites you.

=====

## About the job

See yourself at Radian? We see you here too.

At Radian, we see you. For the person you are and the potential you hold. That's why we've embraced a new way of working that lets our people across the country be themselves, be their best and be their boldest. Because when each of us is truly seen, each of us gives our best — and at Radian, we'll give you our best right back.

Studies have shown that job seekers may hesitate to apply for jobs unless they meet every single qualification listed. We strive to see the potential in each applicant, so if you're excited about this role but your experience or education level doesn't align perfectly with every qualification in the job description, we encourage you to apply anyway. You may be just the right candidate for this or other roles.

#### See Yourself as a Generative AI Engineer

As the Generative AI Engineer you are responsible for the development of innovative AI solutions.

Your expertise in generative models, deep learning, and data analysis will be critical in creating intelligent and transformative AI applications. The resource will work closely with cross-functional teams to conceptualize, design, test, and deploy AI projects that drive innovation and provide value in the rapidly evolving field of artificial intelligence.

See Your Primary Duties And Responsibilities

- Design and develop algorithms for generative models using deep learning techniques.
- Collaborate with cross-functional teams to integrate generative AI solutions into existing workflow systems.
- Research and stay up-to-date on the latest advancements in generative AI technologies and methodologies.
- Optimize and fine-tune generative models for performance and efficiency.
- Troubleshoot and resolve issues related to generative AI models and implementations.
- Create and maintain documentation for generative AI models and their applications.
- Communicate complex technical concepts and findings to non-technical stakeholders.
- Perform other duties as assigned or apparent.

See The Job Specifications

Your Basic Education And Prior Work-Related Experience

- Degree Requirement: Bachelor's Degree or Equivalent Experience
- Work Experience: 1-2 years of prior work-related experience

Additional Qualifications

- Strong Python skills; Effective prompt engineering techniques; familiarity with other languages is a plus.
- Proficiency in AWS Bedrock includes Model access and knowledge base implementations. Experience in Azure OpenAI.
- Hands-on with models for text (e.g., Claude, Openai gpt, Gemini, etc.), understand model fine tuning.
- Building efficient data pipelines and managing large datasets.
- Measuring model outputs with suitable metrics; awareness of bias and fairness.
- AI coding assistant (e.g. github copilot and etc), version control (Git), testing, and building scalable systems.
- Must have AWS serverless experience, Azure or GCP.
- Reading and implementing recent AI papers; experimentation.
- Documenting work and collaborating with teams.

## See Your Location

Radian is committed to a flexible work environment for many of our roles. This is a *\*Work From Anywhere\** role meaning you have the flexibility to work from home (or another designated workspace that fits your needs).

This role provides additional flexibility should you want to work on-site at a Radian office. Explore our office locations [here](#) and let your Talent Acquisition Partner know you would be interested in working on-site.

- Work From Anywhere is subject to Radian's Alternative Work Policy and business needs.

## See Why You Should Work With Us

- Competitive Compensation: anticipated base salary from \$78,700 to \$112,400 based on skills and experience. This position is eligible to participate in an annual incentive program.
- Rest and Relaxation. This role is eligible for 20 days of paid time off annually, which is prorated in the year of hire based on hire date. In addition, based on your hire date, you will be eligible for 9 paid holidays + 2 floating holidays. Parental leave is also offered as an opportunity for all new parents to embrace this exciting change in their lives.
- Our Company Makes an Impact. We've been recognized by multiple organizations like Bloomberg's [Gender-Equality Index](#), HousingWire's [Tech 100](#), and The Forum of Executive Women's [Champion of Board Diversity](#). Radian has also pledged to PwC's [CEO Action for Diversity & Inclusion](#) commitment.
- Comprehensive Health Benefits. Multiple medical plan choices, including HSA and FSA options, dental, vision, and basic life insurance.
- Prepare for your Future. 401(k) with a top of market company match (did we mention the company match is immediately vested?!) and an opportunity to participate in Radian's Employee Stock Purchase Plan (ESPP).
- Homebuyer Perks. Our Homebuyer Perks program helps employees navigate the home searching, buying, selling, and refinancing processes and provides valuable financial benefits to encourage, enable, and support home ownership.
- Additional Benefits. To learn more about our benefits offerings, visit our [Benefits Page](#).

*The application period for the job is estimated to be 20 days from the job posting date. However, this timeline may be shortened or extended depending on business needs and the availability of qualified candidates.*

Radian will consider for employment qualified applicants with arrest or conviction records in a manner consistent with the requirements of the law, including any applicable fair chance law.

## See More About Radian

Radian is a fintech servicing the mortgage and real estate services industry. As a team, we pride ourselves on seeing the potential of every person, every idea and every day.

Seeing each other at Radian goes far beyond our open, flexible culture. It means seeing our people's potential – and creating inspiring career paths that help them get there. Or seeing new pathways and innovating for the future of our industry. It means seeing each other for all that we are. And it means seeing our purpose as one that extends beyond the bottom line – having an impact on communities across the country to help more people achieve the American Dream of homeownership.

We hope you'll see yourself at Radian. See more about us at [Radian.com](https://www.radian.com).

## Defining Roles for Radian's Future

Understanding the qualities and characteristics that define a Leader and an Employee is important to building our future-fit workforce. Radian's future is only as bright as its people. For that reason, our People Plan includes profiles to support the qualities and characteristics that each [Leader](#) as well as each [Employee](#) should embody upon hire or via development.

## EEO Statement

Radian complies with all applicable federal, state, and local laws prohibiting discrimination in employment. All qualified applicants will receive consideration for employment without regard to gender, age, race, color, religious creed, marital status, gender identity, sexual orientation, national origin, ethnicity, ancestry, citizenship, genetic information, disability, protected veteran status or any other characteristic protected by applicable federal, state, or local law.

An applicant's criminal history may have a direct, adverse, and negative relationship with some of the material job duties of this position. The material duties include those listed in the "Primary Duties and Responsibilities" section above, as well as the ability to adhere to Company policies, exercise sound judgment, effectively manage stressful situations, work safely and respectfully with others, exhibit trustworthiness, and safeguard confidential information belonging to the Company and its customers. Pursuant to the California Fair Chance Act, Los Angeles County Fair Chance Ordinance for Employers, Fair Chance Initiative for Hiring Ordinance, and San Francisco Fair Chance Ordinance, we will consider for employment qualified applicants with arrest and conviction records.

## Equal Opportunity Employer Details

To learn more about Radian's Code of Conduct and Ethics and workplace conduct, please click [here](#). Radian participates in E-Verify [Link](#) (en español [Link](#)). Learn more about your rights under immigration laws [Link](#) (en español [Link](#)). View the "Know Your Rights: Workplace Discrimination is Illegal" poster [Link](#). View "Employee Rights under FMLA" [Link](#). View "Employee Rights under EPPA" [Link](#).

=====

## About the job

This is us, your new colleagues!

Our organization is based on people and great teamwork.

We work at the intersection of design, engineering and domain expertise, our passionate team of problem solvers works closely with customers to blaze new trails that will positively impact the future of health. We are a truly global company; we rely on diversity and together we create a workplace that brings the best out of everyone. Using technology and insights, we move nimbly and provide trusted advice, seeking ways to amplify results.

We're committed to bettering healthcare and empowering wellness.

AI Engineer

Remote

Job Description:

Role Overview

You'll be a core member of our AI team, designing, implementing, and optimizing GenAI applications for live, real-time deployments. Our projects frequently involve prompt engineering, vector databases, frameworks like LangChain, Langgraph, and agentic tools such as AutoGen and Crew. You'll help build production-ready solutions for tasks like text generation, summarization, and Q&A systems—often within a cloud environment like Azure, AWS and/or GCP.

Responsibilities

- **Advanced Python Proficiency:** Write clean, efficient, and maintainable Python code for data preprocessing, model training, and integration with production systems.
- **Prompt Engineering:** Design, test, and refine prompts to produce optimal AI-driven outputs (chatbots, content generation, etc.).
- **Retrieval-Augmented Generation (RAG):** Implement techniques that combine model outputs with external data sources to improve accuracy in real-world scenarios.
- **LangChain & Langgraph Integration:** Leverage these libraries to build modular, scalable AI pipelines for language-based tasks.
- **Vector Database Management:** Work with databases like Pinecone, Faiss, or other vector stores to store embeddings and enable efficient semantic search.
- **Chunking Techniques:** Develop methods to split large documents into manageable sections, ensuring more accurate and efficient retrieval and generation.
- **Streamlit (or Similar) Prototyping:** Create user-facing dashboards or internal tools to visualize and interact with AI functionalities.
- **Metrics & Performance Tracking:** Implement monitoring and logging solutions to track performance metrics (response time, accuracy, user engagement) and identify areas of improvement.

- Collaboration & Documentation: Participate in code reviews, maintain thorough documentation, and share best practices with cross-functional teams.

#### What We're Looking For

- Education: Bachelor's degree in Computer Science, Data Science, Engineering, or related field (or equivalent practical experience).
- AI/ML Foundations: Familiarity with NLP, machine learning fundamentals, and deep learning frameworks (e.g., PyTorch or TensorFlow).
- Programming Skills: Solid Python background; comfort with libraries like NumPy, pandas, and familiarity with concurrency or parallelism is a plus.
- Hands-On Project Experience: Coursework or personal projects involving LLMs, vector databases, or text analysis; experience deploying real-time AI solutions is a bonus.
- Cloud Familiarity: Exposure to either Azure or AWS; ability to learn how to set up virtual machines, container services, or serverless functions.
- Adaptability: Excitement about experimenting with new libraries (LangChain, Langgraph) and emerging large language models.
- Curiosity & Team Spirit: Eagerness to learn, collaborate, iterate quickly, and stay current on GenAI trends.

#### Good to Have

- Agentic Frameworks: Experience or interest in frameworks like AutoGen, Crew, or other agentic systems for automating AI-driven tasks.
- MLOps & Tooling: Familiarity with containerization (Docker), CI/CD pipelines, or other DevOps practices for robust AI model deployment.
- Advanced Python Libraries: Proficiency in data processing libraries (e.g., pandas, Dask) and knowledge of parallelization techniques.
- Cloud Deployment: Use either Azure or AWS to configure and deploy AI models, ensuring scalability, security, and optimal performance for live applications.
- Full-Stack Knowledge: Understanding of front-end technologies for better integration of AI-backed features in web or mobile interfaces.
- 

=====

## About the job

### Company Description

Jampack is on a mission to make wholesale radically simple. Our AI agents streamline the end-to-end operations of getting physical products on shelves.

### Role Description

We're building production-grade agents that reason across messy retail systems, autonomously execute workflows, and continually learn from real-world feedback. You'll lead end-to-end — from research to deployment — on capabilities that ship to customers.

### What You'll Do

- Architect and ship agentic systems for ordering, freight tendering, reconciliation, and exception handling
- Build robust tool-use, retrieval, and planning loops with strong safety/guardrails
- Productionize models (fine-tuning, distillation, evals, observability, feedback loops)
- Pair closely with customers to translate workflows into reliable automation

#### Qualifications

- 4+ years building ML/AI products in production (startups preferred)
- Deep experience with LLMs, retrieval, function-calling, and eval frameworks
- Strong software engineering fundamentals (TypeScript/Node, Python, cloud infra)
- Knowledge of software development best practices including version control, testing, and documentation
- Proficiency in data analysis, data engineering, and database management
- Excellent problem-solving abilities and analytical skills
- Ability to work collaboratively in a team environment
- Bias to ship; comfort navigating ambiguity; high ownership

=====

## About the job

### About Anthropic

Anthropic's mission is to create reliable, interpretable, and steerable AI systems. We want AI to be safe and beneficial for our users and for society as a whole. Our team is a quickly growing group of committed researchers, engineers, policy experts, and business leaders working together to build beneficial AI systems.

### About The Role

As an Applied AI team member at Anthropic, you will be a Pre-Sales architect focused on becoming a trusted technical advisor helping large enterprises understand the value of Claude and paint the vision on how they can successfully integrate and deploy Claude into their technology stack. You'll combine your deep technical expertise with customer-facing skills to architect innovative LLM solutions that address complex business challenges while maintaining our high standards for safety and reliability.

Working closely with our Sales, Product, and Engineering teams, you'll guide customers from initial technical discovery through successful deployment. You'll leverage your expertise to help customers understand Claude's capabilities, develop evals, and design scalable architectures that maximize the value of our AI systems.



## Responsibilities

- Partner with account executives to deeply understand customer requirements and translate them into technical solutions, ensuring alignment between business objectives and technical implementation
- Serve as the primary technical advisor to enterprise customers throughout their Claude adoption journey, from discovery to initial evaluation through deployment. You will need to coordinate internally across multiple teams & stakeholders to drive customer success
- Support customers building with both the Claude API and Claude for Work
- Create and deliver compelling technical content tailored to different audiences. You will need to be able to spread the gamut from technical deep dives for engineering & development teams up to business value focused conversations with executives
- Guide technical architecture decisions and help customers integrate Claude effectively into their existing technology stack
- Help customers develop evaluation frameworks to measure Claude's performance for their specific use cases
- Identify common integration patterns and contribute insights back to our Product and Engineering teams
- Travel occasionally to customer sites for workshops, technical deep dives, and relationship building
- Maintain strong knowledge of the latest developments in LLM capabilities and implementation patterns

## You May Be a Good Fit If You Have

- 5+ years of experience in technical customer-facing roles such as Solutions Architect, Sales Engineer, or Technical Account Manager
- Experience working with enterprise customers, navigating complex buying cycles involving multiple stakeholders
- Exceptional ability to build relationships with and communicate technical concepts to diverse stakeholders to include C-suite executives, engineering & IT teams, and more
- Strong technical communication skills with the ability to translate customer requirements between technical and business stakeholders
- Experience designing scalable cloud architectures and integrating with enterprise systems
- Comfortable with python
- Familiarity with common LLM frameworks and tools or a background in machine learning or data science
- Excitement for engaging in cross-organizational collaboration, working through trade-offs, and balancing competing priorities
- A love of teaching, mentoring, and helping others succeed

- Excellent communication and interpersonal skills, able to convey complicated topics in easily understandable terms to a diverse set of external and internal stakeholders. You enjoy engaging in cross-organizational collaboration, working through trade-offs, and balancing competing priorities
- Passion for thinking creatively about how to use technology in a way that is safe and beneficial, and ultimately furthers the goal of advancing safe AI systems

Deadline to apply: None. Applications will be reviewed on a rolling basis.

=====

## About the job

REMOTE!! Start up exp a must:))!

We're looking for a Principal AI Engineer to join our Engineering team and build cutting-edge conversational AI solutions. You'll play a leading role in designing and deploying intelligent chatbots, virtual assistants, recommendation systems, and multi-agent AI frameworks that turn real-world use cases into robust, scalable workflows.

This is a unique opportunity to shape the future of applied AI, working at the forefront of conversational AI innovation in a collaborative, forward-thinking environment.

### Responsibilities

- Build AI/ML pipelines and integrations using Python, LLMs, and LangChain/LlamaIndex.
- Design and implement agent-based AI architectures (Agentic AI) to enable autonomous, multi-step task completion.
- Partner with stakeholders to gather requirements, craft use cases, and translate them into user stories and technical requirements.
- Rapidly prototype and refine AI-driven features with direct client and user feedback.
- Fine-tune large language models for domain-specific tasks, ensuring scalability, accuracy, and efficiency.

### Qualifications

- 3+ years of strong Python development experience.
- Hands-on expertise with LLMs (GPT, LLaMA, etc.).
- Proficiency in LlamaIndex for vector data ingestion, indexing, and retrieval.
- Proven experience building or integrating agentic/multi-agent AI systems.
- Skilled at simplifying complex technical concepts for stakeholders.
- Preferred: background in educational technology, or related fields.

Ability to thrive in a fully remote environment with strong self-management and collaboration skills.

=====

## About the job

### Overview

Spotter is a platform for Creators, providing services and software designed to accelerate growth for

the world's best Creators and brands. Creators working with Spotter can access the capital, knowledge, community, and personalized AI software products they need to succeed. With unique knowledge of how Creators work, the resources they need to grow, and the challenges they face, Spotter is empowering top YouTube Creators to succeed.

Spotter has already deployed over \$950 million to Creators to reinvest in themselves and accelerate their growth, with plans to reach \$1 billion in investment in 2025. With a premium catalog that spans over 725,000 videos, Spotter generates more than 88 billion monthly watch-time minutes, delivering a unique scaled media solution to Advertisers and Ad Agencies that is transparent, efficient, and 100% brand safe. For more information about Spotter, please visit <https://spotter.com>.

This is a remote position and can be based anywhere in the United States.

#### What You'll Do/Responsibilities

As a Senior AI Engineer, you'll design, optimize, and deploy AI-powered conversational and workflow systems that drive personalization, context-awareness, and efficiency for our creator-facing products. You'll work closely with product, engineering, and data science teams to ensure our agentic AI experiences are not only intelligent but deeply relevant and engaging.

- Build, evaluate, and refine agentic AI systems for high-quality chat interactions and workflow automation.
- Develop robust evaluation frameworks to measure performance, accuracy, and user satisfaction.
- Improve the relevance and persistence of conversation memories and contextual grounding.
- Enhance prompts and orchestration logic to improve task completion, personalization, and overall system reliability.
- Experiment with model fine-tuning and domain adaptation when beneficial to user experience.
- Leverage tools such as LangGraph, FalkorDB, and PGVector to manage context and optimize retrieval pipelines.
- Collaborate with product managers, UX designers, and other engineers to align AI capabilities with business goals.

#### Who You Are/Required Qualifications

- 5 years software development experience with at least 1 year of building LLM-focused consumer or enterprise SaaS applications, including designing and optimizing LLM interactions and evaluation workflows.
- Proficiency with Python and modern AI frameworks (LangChain, LangGraph, or similar orchestration tools).
- Experience with vector databases (e.g., FalkorDB, PGVector, Pinecone, Weaviate) and retrieval-augmented generation.
- Knowledge of embedding models, retrieval strategies, and context-window optimization.
- Familiarity with model training, fine-tuning, and evaluation methodologies.
- Understanding of graph databases and their role in storing and retrieving structured knowledge.
- Ability to collaborate in cross-functional, fast-paced environments.

## Why Spotter

- Medical insurance covered up to 100%
- Dental & vision insurance
- 401(k) matching
- Stock options
- Autonomy and upward mobility
- Diverse, equitable, and inclusive culture, where your voice matters.

=====

## About the job

### About Assembled

[Assembled](#) builds the infrastructure that underpins exceptional customer support, empowering companies like CashApp, Etsy, and Robinhood to deliver faster, better service at scale. With solutions for workforce management, BPO collaboration, and AI-powered issue resolution, Assembled simplifies the complexities of modern support operations by uniting in-house, outsourced, and AI-powered agents in a single operating system. Backed by \$70M in funding from NEA, Emergence Capital, and Stripe, and driven by a team of experts passionate about problem-solving, we're at the forefront of support operations technology.

### The Team

Our Agent Intelligence team is focused on maximizing AI agent performance and driving volume automation across voice, chat, and email channels. This team is vital for our vision of creating the ultimate omnichannel AI customer support agent and represents one of the fastest growing areas both inside our company and across the industry. We're currently one of the largest companies deploying AI agents in production at massive scale, working with enterprise clients like Canva, Etsy, and Ramp. The engineering challenges are significant and largely unsolved: building sophisticated

evaluation systems, automating knowledge generation, optimizing AI agent accuracy, and creating the infrastructure needed to continuously improve agent intelligence at scale.

### Your Impact

As part of the Agent Performance Engineering team, you'll be working on the core systems that make our AI agents smarter, more accurate, and more capable of handling complex customer interactions. You'll be building the evaluation frameworks, knowledge automation systems, and intelligence optimization tools that directly impact our ability to automate customer support at enterprise scale.

### Responsibilities

- Build foundational evaluation infrastructure: Develop comprehensive evaluation systems from the ground up, including golden dataset creation, automated benchmarking, and model comparison tools. You'll help create the frameworks that enable us to measure and optimize AI agent performance across all communication channels.
- Automate knowledge generation: Design and implement systems that automatically create synthetic guides, documentation, and metadata to improve agent knowledge bases. You'll work on cutting-edge approaches to knowledge extraction and augmentation.
- Optimize AI agent accuracy: Enhance our retrieval systems, implement advanced prompt optimization techniques, and build tools that continuously improve agent responses through automated evaluation and refinement.
- Develop intelligence infrastructure: Architect systems that enable rapid model upgrades, A/B testing of different AI approaches, and scalable evaluation pipelines that support enterprise deployment.
- Drive volume automation: Focus on the north star goal of maximizing automated resolutions across voice, chat, and email by building the intelligence systems that make it possible.

### About You

You might be a good fit if you:

- Have 5+ years of experience in software engineering as an individual contributor
- Have experience with AI evaluation systems, data pipelines, or AI model optimization
- Are passionate about building systems that measure and improve AI performance
- Have worked with retrieval systems, knowledge bases, or information extraction
- Enjoy building tools and infrastructure that enable other engineers and AI systems to perform better
- Are highly ambitious and driven, setting high goals for yourself and others
- Put customers first, focusing on solving real problems that impact support quality
- Enjoy fast-paced environments and can quickly adjust when new insights emerge
- Have a bit of a maverick streak that helps you come up with creative solutions
- Have made a noticeable impact on small teams and have solid experience in startups or smaller companies
- Stay humble and open to feedback, value teamwork, and are always ready to learn and grow

### Technologies You'll Work With

- Python/Golang for evaluation systems and data processing

- LLMs and prompt optimization frameworks
- Vector databases and retrieval systems
- AI evaluation and benchmarking tools
- Data pipeline and automation infrastructure

The estimated base salary range for this role is \$120,000 - \$270,000 per year. The base pay offered may vary depending on location, job-related knowledge, skills, and experience. Stock options are provided as part of the compensation package, in addition to a full range of medical, financial, and/or other benefits, dependent on the position offered.

Our U.S. benefits

- Generous medical, dental, and vision benefits
- Paid company holidays, sick time, and unlimited time off
- Monthly credits to spend on each: professional development, general wellness, Assembled customers, and commuting
- Paid parental leave
- Hybrid work model with catered lunches everyday (M-F), snacks, and beverages in our SF & NY offices
- 401(k) plan enrollment

## Benefits found in job post

- 401(k)

## About the job

AI Engineer – Agentic Systems & Marketing Workflow Automation

About Inflection.io

inflection.io is reshaping marketing technology with intelligent automation. We empower marketers by leveraging the latest in large language models (LLMs) and AI agentic systems to optimize, automate, and personalize every aspect of the marketing workflow.

### Role Overview

As our AI Engineer, you will lead the end-to-end development of agentic systems that harness LLMs to automate and orchestrate marketing workflows. You'll architect, build, and launch reliable AI-driven components, working closely with technical and product teams, with a direct and measurable impact on our platform and client success.

## Key Responsibilities

- Design, prototype, and implement agentic workflows for marketing automation using LLM frameworks like LangChain.
- Build and maintain multi-agent systems for campaign execution, reporting, content validation, and other automated marketing use cases.
- Integrate LLM-powered features into our SaaS offering, ensuring scalability, security, and data integrity.
- Benchmark and assess LLMs for applicability, performance, and cost in production-grade systems.
- Collaborate cross-functionally to refine requirements and rapidly iterate on workflow and agent capabilities.
- Maintain clear technical documentation and contribute to architectural decisions.

## Required Qualifications

- Bachelor's or Master's degree in Computer Science or equivalent, with emphasis on distributed systems and data management fundamentals.
- Proven experience independently building at least 2 fully functional agentic flows using frameworks such as LangChain, Crew, n8n, or similar in production or live proof-of-concept environments.
- Strong programming experience in Python; proficiency with LLM APIs (OpenAI, Anthropic, etc.) and MCP servers.
- Solid understanding of distributed systems design, data pipelines, and cloud-based service scaling.
- Hands-on experience integrating AI-driven features into real-world SaaS or MarTech platforms.
- Start-up experience, especially in early-stage or high-growth settings, with comfort in ambiguity and fast-paced iteration.

## Nice-to-Haves

- Experience with vector databases (e.g., Pinecone, Weaviate) and RAG pipelines.
- Familiarity with microservices, API design (REST/GraphQL), and DevOps practices.
- Past work orchestrating marketing campaigns or building marketing technology tools.

## Location & Work Policy

- 100% remote; candidates must be based in Seattle or SF Bay Area to maximize collaboration with leadership.

Join inflection.io and drive the next era of AI-powered marketing!

=====

## About the job

### About Meshy

Headquartered in the Silicon Valley, Meshy is the leading 3D generative AI company on a mission to Unleash 3D Creativity. Meshy makes it effortless for both professional artists and hobbyists to create unique 3D assets—turning text and images into stunning 3D models in just minutes. What once took weeks and \$1,000 now takes 2 minutes and \$1.

Our global team of top experts in computer graphics, AI, and art includes alumni from MIT, Stanford, Berkeley, as well as veterans from Nvidia and Microsoft. With 3 million users (and growing), Meshy is trusted by top developers and backed by premiere venture capital firms like Sequoia and GGV.

- No. 1 popularity, among 3D AI tools, according to A16Z games,
- No. 1 website traffic, among 3D AI tools, according to SimilarWeb (2M monthly visits),
- Leading 3D foundation model, delighted texture & fine geometry,
- \$52M funding by Top VCs,
- 2.5M users & 20M models generated!

[Ethan Yuanming Hu](#) serves as the founder and CEO. Ethan got his [Ph.D. in graphics and AI](#) from MIT, where he developed the [Taichi](#) GPU programming language (27K stars on GitHub, used by 300+ institutes). His Ph.D. thesis got a honorable mention of SIGGRAPH 2022 Outstanding Doctoral Dissertation Award and his research has been cited over [2700 times](#). his favorite animal is the llama.

### About The Role

We are seeking research engineers with extensive technical experience in training AI models. You will join our research team to innovate and build our core 3D foundation models, and tackle challenging algorithm design problems.

### What We're Looking For

- Strong mathematical and statistical background with deep insights into machine learning algorithms.
- Familiar with deep learning models (ResNet, Transformer, etc.) and generative AI algorithms (Diffusion, GAN, VAE, Flow-based models, etc.).
- Proficient in using open-source deep learning frameworks like TensorFlow/PyTorch.



- Experience in model training and performance optimization.

Nice to Have :

- 3+ years of practical work experience in AI, with hands-on development experience.
- Experience of multimodal training such as 3D, video generation, or computer vision methods.
- Contributions to popular open-source machine learning projects.
- Publications in top-tier AI/ML/CV conferences.

Our Values

- Brain: We value intelligence and the pursuit of knowledge. Our team is composed of some of the brightest minds in the industry.
- Heart: We care deeply about our work, our users, and each other. Empathy and passion drive us forward.
- Gut: We trust our instincts and are not afraid to take bold risks. Innovation requires courage.
- Taste: We have a keen eye for quality and aesthetics. Our products are not just functional but also beautiful.

Why Join Meshy?

- Competitive salary, equity, and benefits package.
- Opportunity to work with a talented and passionate team at the forefront of AI and 3D technology.
- Flexible work environment, with options for remote and on-site work.
- Opportunities for fast professional growth and development.
- An inclusive culture that values creativity, innovation, and collaboration.
- Unlimited, flexible time off.

Benefits

- Competitive salary, benefits and stock options.
- 401(k) plan for employees.
- Comprehensive health, dental, and vision insurance.
- The latest and best office equipment.

=====

## About the job

AI Agent Developer Evangelist | naptha.ai

About This Role

We are seeking an exceptional AI Agent Developer Evangelist to build and nurture relationships with

frontier AI developers and shape the future of AI agent development. This is a rare opportunity to influence the future of AI agent infrastructure at a massively ambitious scale, backed by industry veterans and technical leaders, NVIDIA Inception, Google for Startups, and Microsoft for Startups.

We're building the foundational infrastructure for the next wave of AI companies, enabling frontier AI developers (many leaving labs like OpenAI, Anthropic, and DeepMind) to build products powered by enormous networks of highly capable next-generation AI agents. As our Developer Evangelist, you'll be the bridge between these technical pioneers and our platform, helping shape both our technology and the broader agent development ecosystem.

### Core Responsibilities

- Build relationships with frontier AI developers and AI infrastructure teams
- Create technical content that showcases advanced agent development patterns
- Develop and share best practices for agent system architecture
- Represent Naptha at technical conferences and AI developer events
- Gather deep technical feedback from the developer community
- Identify emerging patterns and needs in agent development
- Shape our platform's evolution based on developer insights

### Impact Areas

- Building a vibrant community of AI agent developers
- Shaping the technical narrative around agent development
- Creating technical content that pushes the boundaries of what's possible
- Fostering collaboration between frontier AI developers
- Influencing the direction of agent development standards
- Helping developers solve complex agent architecture challenges

You're a good fit if you have:

- Deep technical understanding of AI systems and agent architectures
- Experience building developer communities in frontier tech
- Strong technical writing and speaking abilities
- Track record of creating influential technical content
- Ability to understand and communicate complex technical concepts
- Network within the AI development community
- Passion for developer experience and education

#### Required Experience:

- Hands-on experience building AI systems or developer tools
- Track record of technical community building
- Strong technical communication skills
- Experience with AI frameworks and tools
- Understanding of developer platforms and ecosystems
- Public speaking and technical presentation experience

#### About the hiring process:

- Technical background discussion
- Community strategy presentation
- Content creation exercise
- Team collaboration interview
- Technical communication assessment

#### Compensation & Benefits:

- Highly competitive salary and significant equity stake
- Remote-first work environment
- Full medical, dental, and vision coverage
- Flexible PTO policy
- Conference and travel budget
- Learning and development support

This is a unique opportunity to shape how the next generation of AI developers builds and deploys agent systems. You'll be working at the forefront of AI development, helping create the community and resources that will define the future of agent development.

=====

We are seeking a Software Engineer to join a joint R&D initiative between the Engineering Systems Team and Snowflake AI Research team. This team drives innovation at the intersection of machine learning and developer productivity, with flagship projects such as ML-powered test systems and LLM-powered code review and automation tools. The team supports both product-facing ML initiatives and internal systems optimization, accelerating Snowflake's thought leadership and operational excellence in enterprise AI.

This role is ideal for candidates who have deep hands-on experience in every stage of the software development lifecycle, have a vision on how AI can play a transformative role in it, and a passion for making it happen.

## KEY RESPONSIBILITIES:

- Partner across research and engineering to design, implement, and maintain machine learning systems that directly impact developer workflows and core platform efficiency.
- Build and optimize LLM-based tools for code review, quality automation, and developer assistance, from research prototype to robust production deployment.
- Contribute to the improvement of internal ML and data tooling, pipelines, and collaboration between AI Research and Engineering Systems.
- Operate across the software lifecycle: requirements, design, implementation, delivery, testing, iteration, and support.
- Contribute to a culture of innovation, excellence, and inclusion in a collaborative and fast-paced environment.

## OUR IDEAL SENIOR SOFTWARE ENGINEER WILL HAVE:

- Strong passion for making developers highly productive.
- A track record of technical leadership
- Excellent problem solving skills, and strong CS fundamentals including data structures, algorithms, and distributed systems.
- Fluency in Java, C++, Python or Go.
- Ability to work in a team environment, collaborate well, and mentor junior engineers.
- A strategic mindset and strong sense of what to prioritize for maximum impact.
- Excellent interpersonal communication skills.
- Outstanding written and oral English language skills.
- Strong software engineering fundamentals: code quality, reproducibility, CI/CD best practices, debugging, testing, and documentation.
- Familiarity with ML operations (MLOps), data/feature versioning, and collaborative software development (Git, containers, etc.).
- Strong communication and teamwork skills; able to describe technical tradeoffs and engage both research scientists and software engineers.

## BONUS POINTS FOR EXPERIENCE WITH THE FOLLOWING:

- Developing data or distributed systems internals.
- Familiarity with different types of machine learning approaches (supervised, reinforcement learning, contrastive learning) and their practical constraints in large-scale, dynamic code environments.
- Proven experience with authoring roadmap documents, project plans and RCAs.

- Experience working in or with developer-facing infrastructure, build/test automation, or productionizing research prototypes.
- Experience with large-scale language models (LLMs), agentic/autoML systems, or ML for developer productivity (e.g., code intelligence, automated testing, static analysis).

#### WHY JOIN US:

This is a rare opportunity to collaborate across Snowflake's leading AI research and engineering system organizations, shaping the next generation of intelligent developer tools and ML-powered infrastructure. You'll work with world-class researchers and engineers, operate at scale, and have impact across both the open-source community and internal Snowflake engineering.

The role offers career development opportunities in ML engineering and advanced software systems, in a values-driven, inclusive environment that values curiosity, impact, and technical excellence.

*Every Snowflake employee is expected to follow the company's confidentiality and security standards for handling sensitive data. Snowflake employees must abide by the company's data security plan as an essential part of their duties. It is every employee's duty to keep customer information secure and confidential.*

Snowflake is growing fast, and we're scaling our team to help enable and accelerate our growth. We are looking for people who share our values, challenge ordinary thinking, and push the pace of innovation while building a future for themselves and Snowflake.

How do you want to make your impact?

For jobs located in the United States, please visit the job posting on the Snowflake Careers Site for salary and benefits information: [careers.snowflake.com](https://careers.snowflake.com)

The following represents the expected range of compensation for this role:

- The estimated base salary range for this role is \$222,000 - \$339,200.
- Additionally, this role is eligible to participate in Snowflake's bonus and equity plan.

The successful candidate's starting salary will be determined based on permissible, non-discriminatory factors such as skills, experience, and geographic location. This role is also eligible for a competitive [benefits](#) package that includes: medical, dental, vision, life, and disability insurance; 401(k) retirement plan; flexible spending & health savings account; at least 12 paid holidays; paid time off; parental leave; employee assistance program; and other company benefits.

=====

## About the job

AI/ML Engineer – Agentic AI, LLM Fine-Tuning & Graph Analytics

### Company Overview

At Flowmentum, we partner with leading organizations on their most ambitious AI initiatives – from building agentic AI systems for enterprise QA automation to tackling real-world challenges in national security, cybersecurity, and cloud innovation.

We're a mission-driven, remote-first team passionate about transforming research into production systems. Our work spans agentic AI strategy, LLM orchestration, graph analytics, and scalable cloud deployments – ensuring AI is secure, reliable, and impactful.

### Role Overview

We're looking for an AI/ML Engineer who thrives at the intersection of research and engineering. In this role, you'll:

- Fine-tune and deploy large language models (LLMs) across distributed systems and cloud platforms (AWS Bedrock, SageMaker, ECS).
- Build and optimize Retrieval-Augmented Generation (RAG) pipelines for structured + unstructured data.
- Design and implement agentic AI workflows that automate complex QA and testing environments.
- Apply graph analytics for anomaly detection, cybersecurity, and systems modeling.
- Contribute to full-stack AI products using TypeScript, React, Next.js.
- Develop scalable ingestion pipelines (PostgreSQL, Apache NiFi, Elasticsearch) to support AI/ML workflows.

### What We're Looking For

- 2+ years of hands-on experience in ML, NLP, or AI research/engineering (internships & research count).
- Strong experience fine-tuning LLMs with PyTorch, Hugging Face, or JAX.
- Proficiency in Python + TypeScript, with experience in React/Next.js.
- Knowledge of distributed systems and cloud AI platforms (AWS preferred).
- Solid SQL + data modeling experience; exposure to graph databases (Neo4j) is a plus.
- Bonus: familiarity with Docker, Slurm, or Azure AI Services.


### Soft Skills

- Independent, proactive problem solver who thrives in remote-first teams.

- Strong communicator with a passion for learning and applying new AI/ML techniques.

#### What We Offer

- Competitive salary + performance bonus.
- Remote-first culture with flexible hours.
- Opportunity to shape next-generation agentic AI systems with real-world impact.

 Join us at Flowmentum and help build the future of secure, scalable AI systems.

=====

## About the job

Cresta is on a mission to turn every customer conversation into a competitive advantage by unlocking the true potential of the contact center. Our platform combines the best of AI and human intelligence to help contact centers discover customer insights and behavioral best practices, automate conversations and inefficient processes, and empower every team member to work smarter and faster. Born from the prestigious Stanford AI lab, Cresta's co-founder and chairman is [Sebastian Thrun](#), the genius behind Google X, Waymo, Udacity, and more. Our leadership also includes CEO, [Ping Wu](#), the co-founder of Google Contact Center AI and Vertex AI platform, [and CTO & co-founder, Tim Shi](#), an early member of Open AI.

We've assembled a world-class team of AI and ML experts, go-to-market leaders, and top-tier investors including Andreessen Horowitz, Greylock Partners, Sequoia, and former AT&T CEO John Donovan. Our valued customers include brands like Intuit, Cox Communications, Hilton, and Carmax and we've been recognized by Forbes and Bain Consulting as one of the top private AI companies in the world.

Join us on this thrilling journey to revolutionize the workforce with AI. The future of work is here, and it's at Cresta.

#### About The Role:

As a Conversation Designer, you will be responsible for ensuring best-in-class high quality AI conversation models that offer real-time coaching and assistance for our customers' contact center agents, and designing conversational experiences for virtual agents for both voice and chat. The role is highly collaborative and will require that you work closely with cross-functional team members from Machine Learning, Data Science, Project Management, Customer Success, Product Leadership,

and Engineering. You will also be responsible for working directly with customers to get alignment on our overall solution design and to enable customers to build their own models using Cresta Opera, our proprietary self-service tool.

#### Responsibilities:

- Lead customer AI Delivery projects, designing, building, deploying, and tuning AI solutions for multiple customers simultaneously
- Collaborate with internal and external stakeholders to design and build solutions that help customers achieve their business objectives using Cresta's AI product suite
- Design solutions for contact centers in sales, retention, collections, and customer service across a range of verticals, from financial services to telecommunications and everything in between
- Lead model maintenance efforts to ensure that Cresta's AI models continue to meet our quality standards
- Deliver Cresta's Opera Enablement Program to key customer stakeholders, empowering them to design, build, and QA their own AI Models using Cresta Opera
- Design and develop virtual agent solutions for customers using flow-based, prompt-based, and hybrid virtual agent platforms

#### Qualifications We Value:

- Strong communication skills (verbal and written)
- Ability to appropriately prioritize and manage multiple projects simultaneously
- Desire to work in a fast-paced environment with the opportunity to learn and develop a wide range of skills
- 2 years of experience in conversational design as a Conversation Designer/Architect (e.g. building bots, conversational AI)
- Experience presenting complex solutions to key customer stakeholders
- Fluency in a second language (preferred)
- Background in Linguistics preferred
- Experience with prompt engineering preferred
- Experience with Dialogflow CX preferred

#### Perks & Benefits:

- We offer Cresta employees a variety of medical, dental, and vision plans, designed to fit you and your family's needs
- Paid parental leave to support you and your family
- Monthly Health & Wellness allowance
- Work from home office stipend to help you succeed in a remote environment



- Lunch reimbursement for in-office employees
- PTO: Flexible

## About the job

Are you a proactive problem-solver with a passion for unlocking the power of AI? We're hiring an AI Prompt Engineer to scale how teams across the company adopt and integrate AI into their day-to-day work. This isn't just about writing prompts—it's about shaping the future of work at our company. You'll be at the forefront of our AI strategy, designing and implementing workflows that drive tangible business outcomes, from improving team efficiency to delivering faster, more consistent customer experiences. You'll build and maintain a centralized prompt library that powers these workflows, ensuring every interaction is clear, reliable, and aligned with our standards. This is your opportunity to directly impact how we operate and empower our teams to achieve more.

### On a Typical Day, You Might

- Design and implement AI-powered workflows in collaboration with internal teams to drive measurable business outcomes.
- Build and maintain a centralized prompt library with high-quality, reusable prompts.
- Iterate on prompts to enhance clarity, consistency, and automation logic.
- Troubleshoot and resolve issues related to prompt execution or system integration.
- Collaborate cross-functionally with Operations, Quality, and leadership to scale review processes and align on priorities.
- Improve the accessibility and structure of internal knowledge content used by AI systems.
- Analyze usage data and monitor workflow performance to generate insights, inform decisions, and identify opportunities for continuous improvement.
- Translate feedback into actionable prompt or workflow improvements.

### The Ideal Candidate Will Bring

- Strong practical understanding of LLMs and core AI concepts (e.g., prompt engineering, RAG, function calling) and how to apply them to internal operations.
- Proven experience in operations or AI-enabled roles, with a focus on building automations that leverage AI to drive meaningful outcomes.
- Familiarity with workflow tools like Zapier, APIs, or internal scripting frameworks to build lightweight AI integrations.
- Great problem-solving skills with an ability to identify, diagnose, and resolve technical issues in AI-powered systems.
- A bias for action, with a passion for moving from idea to implementation and thriving in environments where experimentation leads to meaningful impact.

- Strong communication and collaboration skills, adept at working cross-functionally with internal teams and external vendors to achieve shared goals.
- Agility in adapting to new technologies, shifting priorities, and evolving business needs.

## About ActiveCampaign

ActiveCampaign is an AI-first, end-to-end marketing platform for people at the heart of the action. It empowers teams to automate their campaigns with AI agents that imagine, activate, and validate—freeing them from step-by-step workflows and unlocking limitless ways to orchestrate their marketing.

With AI, goal-based automation, and 950+ app integrations, agencies, marketers, and owners can build cross-channel campaigns in minutes—fine-tuned with billions of data points to drive real results for their unique business.

ActiveCampaign is the trusted choice to help businesses unlock a new world of boundless opportunities—where ideas become impact and potential turns into real results.

As a global multicultural company, we are proud of our inclusive culture which embraces diverse voices, backgrounds, and perspectives. We don't just celebrate our differences, we believe our diversity is what empowers our innovation and success. You can find out more about our DEI initiatives [here](#).

## Perks And Benefits

At ActiveCampaign, we prioritize employees' well-being and professional growth by cultivating a culture centered on collaboration and innovation. When you join our team, you'll not only have the opportunity to make a significant impact, but also enjoy a range of benefits tailored to support your personal and career development.

## Here Are Some Of The Benefits We Offer

- **Comprehensive Health & Wellness:** Top-tier benefits package that includes a fully-covered High Deductible Health Plan (HDHP), complimentary access to telehealth services, and a free subscription to Calm.
- **Growth & Development:** Access to LinkedIn Learning, professional development programs, and career growth opportunities in a fast-growing organization.

- Generous Paid Time Off: Recharge and take the time you need to maintain work-life balance with open PTO.
- Total Rewards: Generous 401(k) matching with immediate vesting, quarterly perks with commuter and lunch benefits for hub based employees or a stipend for remote workers, and a four-week paid sabbatical with bonus after five years.
- Collaborative Culture: Work alongside brilliant, passionate colleagues in an environment that values innovation, teamwork, and mutual support.

ActiveCampaign is an equal opportunity employer. We recruit, hire, pay, grow and promote no matter of gender, race, color, sexual orientation, religion, age, protected veteran status, physical and mental abilities, or any other identities protected by law.

Our Employee Resource Groups (ERGs) strive to foster a diverse inclusive environment by supporting each other, building a strong sense of belonging, and creating opportunities for mentorship and professional growth for their members.

=====

## About the job

### About Us and the Role

At Sobek AI, we are building the secure nervous system for the next generation of life-science innovation networks and government collaborations. We're backed by \$10M+ in grants and funding, and work with global, high-impact partners such as PATH and Institut Pasteur de Dakar.

We're looking for a bedrock hire to help us build out the core that makes this possible.

You will build and own foundational services that every Sobek product relies on: Secure, high-throughput event ingestion; policy-aware orchestration; end-to-end, reliable LLM integration and first-class observability. The platform you build will directly power our mission to accelerate critical Vaccine R&D and global emergency coordination.

This is a high-impact, high-ownership role for someone who is obsessed with building reusable, enterprise-grade services. You won't just be using AI; you'll also be building the hardened, scalable infrastructure that makes AI trustworthy for the world's most sensitive data, with a direct line to better public-health outcomes.

### What You'll Do

#### Architect Core Platform Functionality

- Design and ship distributed systems including continuous event-ingestion, and task-processing pipelines (e.g. AWS SQS, Lambda functions).
- Expose clean and highly available REST/gRPC APIs that power real-time permissioning decisions and agent orchestration.

#### Engineer Production-Grade AI Systems

- Own orchestration engines that reliably manage complex workflows across multiple LLMs and other AI services, building the intelligent control plane for cost, latency, and security.
- Solve the unique systems challenges of production AI: engineering for non-determinism, guaranteeing data provenance throughout complex transformations, and building the operational hooks needed to monitor, debug, and continuously improve our AI systems.

#### Instill Security and Reliability by Design

- Bake enterprise-grade security into every service, implementing multi-tenant isolation, fine-grained IAM, encryption-at-rest, and auditable trails that meet stringent compliance needs.
- Own the operational excellence of your services by implementing and monitoring comprehensive instrumentation using tools like Prometheus and OpenTelemetry.

#### Lead Through Engineering Excellence

- Serve as a pillar of engineering quality, writing clear Architecture Decision Records and providing thoughtful code reviews.

#### About You

You are a systems thinker at heart, energized by the challenge of building elegant, resilient platforms that solve novel problems. You believe that great AI infrastructure is a product in itself and a force multiplier for the entire company.

- You have 4+ years of experience shipping complex backend systems in high-growth environments where quality and speed are not mutually exclusive.
- You thrive at the intersection of distributed systems and machine learning. You see LLMs not just as a black box to call, but as a new class of architectural component that requires first-principles thinking to deploy safely and reliably at scale.
- You possess a first-principles understanding of distributed systems: Concurrency, consistency, and event-driven architectures are your native language.
- You have expert-level fluency in Python or TypeScript/Node.js, and you're comfortable operating across synchronous APIs and asynchronous worker pools.
- You have a hands-on cloud background and experience turning architectural diagrams into reality with IaC and container orchestration.
- You have a profound sense of ownership, taking pride in the operational excellence of your services and instinctively mentoring others through the quality of your work.
- You are a clear communicator and an empathetic collaborator.

#### Details

- Compensation: \$140 K – \$190 K + equity
- Location: Hybrid (Seattle, WA)
- Visa: We do not sponsor visas for this role at this time
- Benefits: Company-paid health coverage (including dependents)

=====

## About the job

## Job Description: Prompt Engineer – Sales Simulations

Position: Prompt Engineer

Type: Full-time

Location: Remote (with potential client interactions)

### Role Overview

We are seeking a Prompt Engineer with expertise in designing, testing, and refining chatbots tailored for sales training simulations. The ideal candidate has experience with complex prompts for AI systems, understands sales processes, and can translate business objectives into impactful AI-driven solutions. Fluency in English is essential, with proficiency in additional languages considered a strong advantage.

### Key Responsibilities

- Prompt Development:
  - Design and optimize chatbot prompts that simulate realistic sales conversations aligned with client-specific sales processes and objectives.
  - Integrate client training goals and evaluation criteria into chatbot interactions to ensure meaningful and actionable feedback.
- Customization & Feedback:
  - Adapt prompts to reflect industry-specific scenarios, focusing on healthcare sectors.
  - Analyze chatbot performance and refine prompts based on feedback from client evaluations.
- Cross-functional Collaboration:
  - Work closely with clients, trainers, and subject matter experts to understand their requirements and align simulations with their training goals.
  - Partner with the technical and product teams to ensure seamless integration and functionality of the chatbot.
- Innovation & Agility:
  - Stay updated on advancements in AI and prompt engineering to continuously improve and innovate our sales training solutions.
  - Demonstrate agility in learning and applying new techniques to deliver high-quality results.

## Required Qualifications

- Experience: Proven experience in designing effective prompts for AI chat systems, preferably in a sales training or customer interaction context.
- Languages: Fluency in English is mandatory; proficiency in other languages is a strong plus.
- Sales Knowledge: Familiarity with sales processes, training objectives, and evaluation frameworks.
- Problem-solving: Strong analytical skills to evaluate prompt performance and adapt strategies to meet client needs.
- Communication: Excellent communication skills to liaise with clients and internal teams effectively.

If you are passionate about AI, have a knack for sales process alignment, and excel at creating engaging chatbot experiences, we would love to hear from you!

=====

## About the job

### About the Role

We are seeking a highly skilled AI Engineer to design, build, and deploy custom AI agents on top of OpenWebUI and other enterprise AI platforms. This role will work directly with our AI Architect and client stakeholders to create scalable, production-ready agent solutions that extend the capabilities of LLM-powered applications.

The ideal candidate has a background in software engineering with strong hands-on experience in Python, containerized environments (Docker, Kubernetes), and integrating APIs. You will play a central role in shaping next generation of intelligent tools, enabling consulting teams and operations groups to leverage AI effectively.

### Responsibilities

- Develop, test, and deploy custom AI agents using OpenWebUI and related frameworks.
- Integrate external APIs, knowledge bases, and business systems into agent workflows.
- Collaborate with the AI Architect and client teams to translate business use cases into technical agent implementations.
- Monitor, optimize, and debug agent performance in production environments.
- Build reusable frameworks, modules, and templates to accelerate future agent development.
- Write documentation and provide knowledge transfer to client engineering teams.
- Stay up to date with emerging LLM architectures, prompt engineering techniques, and agent orchestration methods.

## Qualifications

- Bachelor's degree in Computer Science, Engineering, or related field (or equivalent experience).
- 3+ years of hands-on software engineering experience (Python strongly preferred).
- Proven experience building applications with LLMs (e.g., OpenAI, Anthropic, Llama).
- Strong understanding of REST APIs, vector databases, and RAG (retrieval-augmented generation) techniques.
- Experience with Docker and cloud environments (Azure or AWS preferred).
- Familiarity with frontend frameworks (React, Streamlit, or similar) is a plus.
- Strong debugging, problem-solving, and collaboration skills.
- Excellent written and verbal communication skills.

## Preferred Skills

- Prior experience with OpenWebUI or LangChain-based frameworks.
- Knowledge of enterprise security, data governance, and compliance considerations in AI systems.
- Experience building internal tools, chatbots, or workflow automation solutions.

## Why Join Us?

- Work on cutting-edge AI adoption projects for enterprise clients.
- Opportunity to design real-world, production-ready AI agents at scale.
- Collaborative environment with an emphasis on experimentation, iteration, and practical results.
- Competitive compensation with clear career growth in AI engineering and architecture.